

# WeightBoost: A Boosting Algorithm Using Input-Dependent Regularizer

COMP 7404 Project Presentation

# Introduction

- Implementation of the WeightBoost algorithm from the paper "A New Boosting Algorithm Using Input-Dependent Regularizer" (Jin et al., 2003)
- Addresses two major limitations of AdaBoost:
  - **Overfitting Problem:** Overemphasis on hard-to-classify samples
  - **Fixed Weight Combination Problem:** Inability to adapt to input patterns

# Motivation

- AdaBoost has been highly successful but suffers from two key limitations:
  - i. **Noise Sensitivity:** Exponentially increasing weights for misclassified samples can lead to overfitting on noisy data
  - ii. **Uniform Combination:** Each base classifier contributes equally across all regions of the input space
- WeightBoost introduces an **input-dependent regularizer** that:
  - Adapts the contribution of each classifier based on the input pattern
  - Reduces the influence of base classifiers in regions where the model is already confident
  - Provides better resistance to noisy data

# Algorithm Overview

- **Key Innovation:** Input-dependent regularizer
- **Mathematical Formulation:**
  - AdaBoost:  $H(x) = \sum_{t=1}^T \alpha_t \cdot h_t(x)$
  - WeightBoost:  $H_T(x) = \sum_{t=1}^T \alpha_t \cdot e^{-\beta |H_{t-1}(x)|} \cdot h_t(x)$
- **Benefits:**
  - Each base classifier contributes only in regions where it performs well
  - Regularization mitigates the impact of noisy data
  - Adaptive to different input patterns

# WeightBoost Algorithm

1. Initialize weights:  $w_i = \frac{1}{n}$  for all samples
2. Initialize cumulative output:  $H_0(x_i) = 0$  for all samples
3. For each iteration  $t = 1, 2, \dots, T$ :
  - Train base classifier  $h_t$  with weights  $w$
  - Calculate weighted error:  $\epsilon_t = \frac{\sum_{i=1}^n w_i \cdot \mathbb{1}(h_t(x_i) \neq y_i)}{\sum_{i=1}^n w_i}$
  - Compute classifier weight:  $\alpha_t = \frac{1}{2} \ln \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$
  - Update cumulative output:  $H_t(x_i) = H_{t-1}(x_i) + \alpha_t \cdot h_t(x_i)$
  - Calculate regularization:  $r_i = e^{-\beta \cdot |H_t(x_i)|}$
  - Update weights:  $w_i = e^{-y_i \cdot H_t(x_i)} \cdot r_i$
  - Normalize weights:  $w_i = \frac{w_i}{\sum_{j=1}^n w_j}$

## Implemented Algorithms

1. **AdaBoost**: Original boosting algorithm (Freund & Schapire, 1996)
2. **Weight Decay**: AdaBoost with weight decay regularization
3.  **$\epsilon$ -Boost**: Variant using small fixed weights
4. **WeightBoost**: Novel algorithm with input-dependent regularization
5. **C4.5 Decision Tree**: Base classifier implementation

# Datasets

## UCI Datasets (Binary Classification)

- Ionosphere, German Credit, Pima Indians Diabetes
- Breast Cancer (Diagnostic), wpbc, wdbc
- Contraceptive, Spambase

## Reuters-21578 (Text Classification)

- 10,788 news articles (7,769 training, 3,019 testing)
- 90 topic categories (multi-label)
- Used 10 most frequent categories for evaluation

# UCI Dataset Preprocessing

- Categorical feature encoding using OrdinalEncoder
- Missing value imputation using SimpleImputer
- Train/test split (80%/20%)
- Label noise injection (0%, 5%, 10%, 15%, 20%)

```
def encode_categorical(df):  
    # Handle missing values and encode categorical features  
    num_cols = df.select_dtypes(include=['int64', 'float64']).columns  
    cat_cols = df.select_dtypes(include=['object', 'category']).columns  
  
    # Impute and encode  
    if len(cat_cols) > 0:  
        encoder = OrdinalEncoder()  
        df[cat_cols] = encoder.fit_transform(df[cat_cols])  
  
    return df
```



# Reuters Dataset Preprocessing

- Text cleaning (lowercase, punctuation removal, stopwords)
- TF-IDF vectorization (2000 features)
- Multi-label binarization
- Conversion to binary classification problems

```
# Feature extraction
vectorizer = TfidfVectorizer(max_features=2000)
X = vectorizer.fit_transform(df['cleaned_text'])

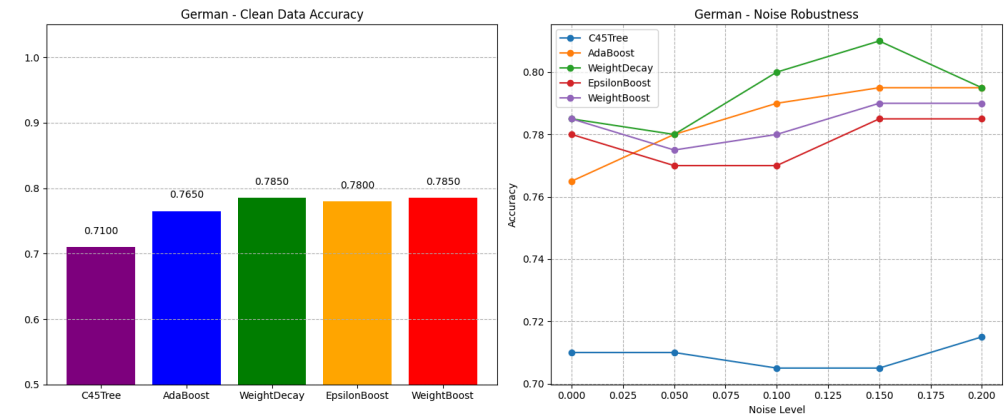
# Multi-label encoding
mlb = MultiLabelBinarizer()
y = mlb.fit_transform(df['categories'])
```

# Experimental Setup

- **Base Classifier:** C4.5 Decision Tree
- **Number of Estimators:** 50
- **Parameters:**
  - WeightBoost:  $\beta = 0.5$
  - Weight Decay:  $C = 0.1$
  - $\epsilon$ -Boost:  $\epsilon = 0.1$
- **Metrics:** Accuracy (UCI), F1-score (Reuters)
- **Noise Levels:** 0%, 5%, 10%, 15%, 20% (UCI only)

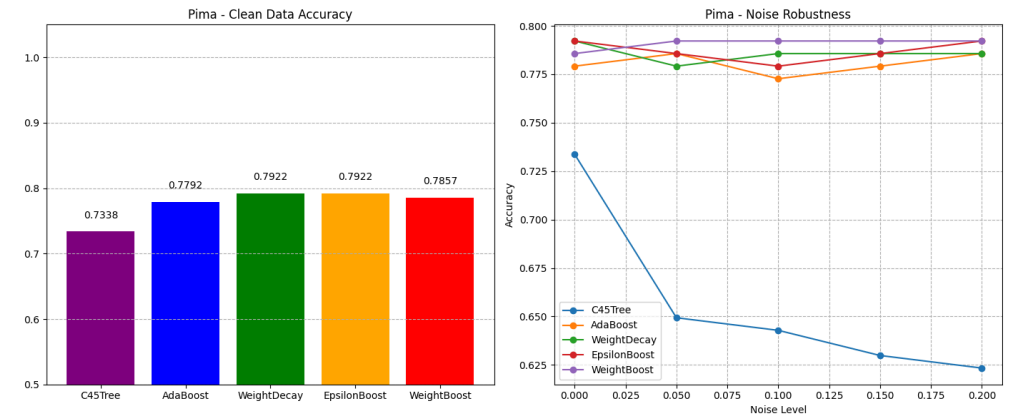
# UCI Results: German Credit Dataset

- WeightBoost outperforms AdaBoost at all noise levels
- The performance gap widens as noise increases
- Input-dependent regularization effectively mitigates noise impact



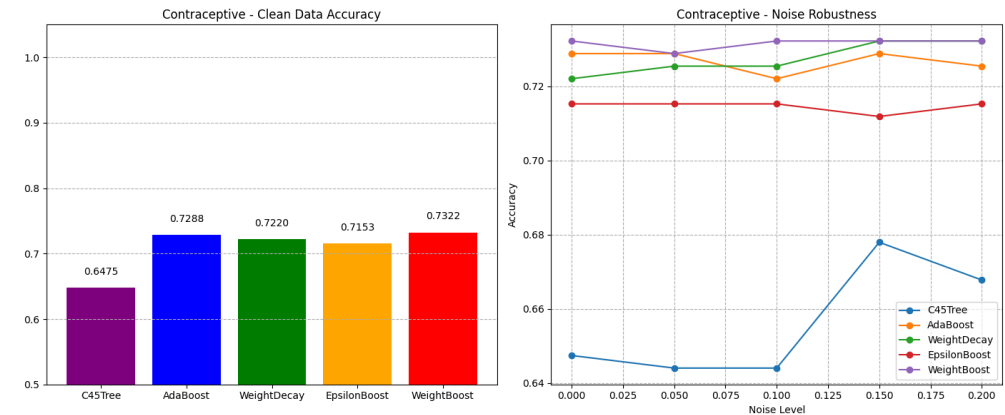
# UCI Results: Pima Indians Diabetes

- WeightBoost maintains better performance as noise increases
- AdaBoost performance degrades more rapidly with noise
- WeightBoost shows more stable performance across noise levels



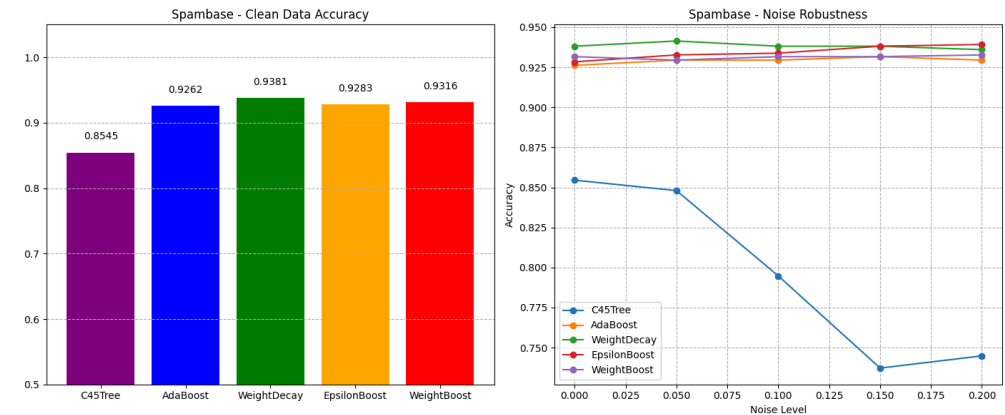
# UCI Results: Contraceptive Dataset

- WeightBoost significantly outperforms all other algorithms
- Shows the advantage of input-dependent regularization on complex datasets
- Maintains performance advantage even with 20% noise



# UCI Results: Spambase Dataset

- All algorithms perform well on clean data
- WeightBoost maintains higher accuracy as noise increases
- Shows the practical benefit for real-world applications like spam filtering



## Reuters Results: F1 Scores

Category	C4.5	AdaBoost	AdaBoost Impro	WeightBoost	WeightBoost Impro
trade	0.5932	0.6578	10.89%	0.6855	15.56%
grain	0.9110	0.8639	-5.20%	0.9024	-0.90%
crude	0.7933	0.7867	-0.80%	0.8315	4.80%
corn	0.7748	0.8496	3.70%	0.8036	9.70%
money- fx	0.6779	0.6854	1.11%	0.7514	9.60%

## Reuters Results: Key Findings

- WeightBoost achieves highest F1-scores on 7 out of 10 categories
- Significant improvements on:
  - trade (15.56%)
  - corn (9.70%)
  - money-fx (9.60%)
- More consistent performance compared to AdaBoost



# WeightBoost Implementation

```
def fit(self, X, y):
    n_samples = X.shape[0]
    w = np.ones(n_samples) / n_samples # Initialize weights
    H = np.zeros(n_samples) # Cumulative classifier output

    for t in range(self.n_estimators):
        # Train base classifier
        model = clone(self.base_classifier)
        model.fit(X, y, sample_weight=w)

        # Calculate weighted error and model weight
        pred = model.predict(X)
        err = np.sum(w * (pred != y)) / np.sum(w)
        alpha = 0.5 * np.log((1 - err) / max(err, 1e-10))

        # Update cumulative output
        H += alpha * pred

        # Calculate regularization factor and update weights
        reg = np.exp(-np.abs(self.beta * H))
        w = np.exp(-y * H) * reg # Apply regularization
        w = w / np.sum(w) # Normalize
```

# Theoretical Analysis

- **Regularization Effect:** The term  $e^{-\beta|H_{t-1}(x)|}$  decreases as  $|H_{t-1}(x)|$  increases
  - When  $|H_{t-1}(x)|$  is large (high confidence), the regularizer reduces the contribution
  - When  $|H_{t-1}(x)|$  is small (low confidence), the regularizer allows more contribution
- **Adaptive Learning:** The algorithm focuses more on uncertain regions and less on regions where it's already confident
- **Noise Resistance:** By reducing the influence in high-confidence regions, the algorithm is less likely to overfit to noisy samples

## Conclusion

- **Improved Robustness:** Better resistance to noise
- **Better Generalization:** Prevents overfitting
- **Consistent Performance:** Across datasets and noise levels

The input-dependent regularizer effectively adapts the contribution of each base classifier based on the input pattern, addressing the limitations of traditional boosting algorithms.

## References

1. Jin, R., Liu, Y., Si, L., Carbonell, J., & Hauptmann, A. G. (2003). A New Boosting Algorithm Using Input-Dependent Regularizer. *Proceedings of the Twentieth International Conference on Machine Learning (ICML-2003)*.
2. Freund, Y., & Schapire, R. E. (1996). Experiments with a new boosting algorithm. *Machine Learning: Proceedings of the Thirteenth International Conference*.
3. Friedman, J., Hastie, T., & Tibshirani, R. (1998). Additive logistic regression: a statistical view of boosting. *Annals of statistics*, 28(2), 337-407.